

JW6 Labs 2 and 3

Rayan Hassan

Instructions and time-saving hints. The purpose of this lab is to learn linear regression in R. You will reproduce the commands blocks in JW Section 6.5 in this Rmarkdown file. **Please do not retype all the R commands as shown in Section 6.5.** Instead, go to link to text, find Chapter 6 Lab text, and cut and paste the text into your Rmarkdown file. Then break up the R commands into the R chunks shown in the text. **Do not put all commands into the same block!**

After breaking the text into the correct chunks, you may knit the document. You will notice some problems when knitting. This is due to the fact that the pasted commands are meant for the command line, not Rmarkdown. You must find Rmarkdown conversions for these problems. One problem is the appearance of the `fix()` commands in the blocks, which displays dataframes in a different window. As a substitute for `fix()`, please use the `kable` command from library `knitr`, as in `knitr::kable()`. The results, however, will be too long to hand in! Use this only for your personal use. For your homework/lab submission, please instead use `head()` instead of `fix()`. The TA will be less burdened that way.

Since R chunks are independent pieces of code, you will sometimes produce errors when imitating the format in the book. For example, the `abline()` commands must be grouped in the same chunk as the initial `plot()` command.

Congratulations! When the document is successfully knitted, you should read Section 6.5 along with the Rmarkdown output. You should see the commands along with the output plots. Together, JW 6.5 and your Rmarkdown file will help you to learn linear regression in R. Note that the plotting is done with core R commands and not with the `ggplot2` package. Submit your Rmd file along with an unzipped PDF of the result before the deadline.

Chapter 6 Lab 2: Ridge Regression and the Lasso

```
x=model.matrix(Salary~.,Hitters)[-1]
y=Hitters$Salary
```

Ridge Regression

```
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 4.3.2
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-8
```

```
grid=10^seq(10,-2,length=100)
ridge.mod=glmnet(x,y,alpha=0,lambda=grid)
```

```
dim(coef(ridge.mod))
```

```
## [1] 20 100
```

```
ridge.mod$lambda[50]
```

```
## [1] 11497.57
```

```
coef(ridge.mod)[,50]
```

```
## (Intercept)      AtBat      Hits      HmRun      Runs
## 407.356050200  0.036957182  0.138180344  0.524629976  0.230701523
##      RBI      Walks      Years      CAtBat      CHits
## 0.239841459  0.289618741  1.107702929  0.003131815  0.011653637
##      CHmRun      CRuns      CRBI      CWalks      LeagueN
## 0.087545670  0.023379882  0.024138320  0.025015421  0.085028114
##      DivisionW      PutOuts      Assists      Errors      NewLeagueN
## -6.215440973  0.016482577  0.002612988 -0.020502690  0.301433531
```

```
sqrt(sum(coef(ridge.mod)[-1,50]^2))
```

```
## [1] 6.360612
```

```
ridge.mod$lambda[60]
```

```
## [1] 705.4802
```

```
coef(ridge.mod)[,60]
```

```
## (Intercept)      AtBat      Hits      HmRun      Runs      RBI
## 54.32519950  0.11211115  0.65622409  1.17980910  0.93769713  0.84718546
##      Walks      Years      CAtBat      CHits      CHmRun      CRuns
## 1.31987948  2.59640425  0.01083413  0.04674557  0.33777318  0.09355528
##      CRBI      CWalks      LeagueN      DivisionW      PutOuts      Assists
## 0.09780402  0.07189612  13.68370191 -54.65877750  0.11852289  0.01606037
##      Errors      NewLeagueN
## -0.70358655  8.61181213
```

```
sqrt(sum(coef(ridge.mod)[-1,60]^2))
```

```
## [1] 57.11001
```

```
predict(ridge.mod,s=50,type="coefficients")[1:20,]
```

```
##      (Intercept)      AtBat      Hits      HmRun      Runs
##  4.876610e+01 -3.580999e-01  1.969359e+00 -1.278248e+00  1.145892e+00
##      RBI      Walks      Years      CAtBat      CHits
##  8.038292e-01  2.716186e+00 -6.218319e+00  5.447837e-03  1.064895e-01
##      CHmRun      CRuns      CRBI      CWalks      LeagueN
##  6.244860e-01  2.214985e-01  2.186914e-01 -1.500245e-01  4.592589e+01
##      DivisionW      PutOuts      Assists      Errors      NewLeagueN
## -1.182011e+02  2.502322e-01  1.215665e-01 -3.278600e+00 -9.496680e+00
```

```
set.seed(1)
train=sample(1:nrow(x), nrow(x)/2)
test=(-train)
y.test=y[test]
```

```
ridge.mod=glmnet(x[train,],y[train],alpha=0,lambda=grid, thresh=1e-12)
ridge.pred=predict(ridge.mod,s=4,newx=x[test,])
mean((ridge.pred-y.test)^2)
```

```
## [1] 142199.2
```

```
mean((mean(y[train])-y.test)^2)
```

```
## [1] 224669.9
```

```
ridge.pred=predict(ridge.mod,s=1e10,newx=x[test,])
mean((ridge.pred-y.test)^2)
```

```
## [1] 224669.8
```

```
ridge.pred=predict(ridge.mod,s=0,newx=x[test,],exact=T,x=x[train,],y=y[train])
mean((ridge.pred-y.test)^2)
```

```
## [1] 168588.6
```


```
lm(y~x, subset=train)
```

```
##
## Call:
## lm(formula = y ~ x, subset = train)
##
## Coefficients:
## (Intercept)      xAtBat      xHits      xHmRun      xRuns      xRBI
##    274.0145     -0.3521     -1.6377      5.8145      1.5424      1.1243
##      xWalks      xYears      xCAtBat      xCHits      xCHmRun      xCRuns
##      3.7287     -16.3773     -0.6412      3.1632      3.4008     -0.9739
##      xCRBI      xCWalks      xLeagueN      xDivisionW      xPutOuts      xAssists
##     -0.6005      0.3379     119.1486     -144.0831      0.1976      0.6804
##      xErrors      xNewLeagueN
##     -4.7128     -71.0951
```

```
predict(ridge.mod,s=0,exact=T,type="coefficients",x=x[train,],y=y[train])[1:20,]
```

```
## (Intercept)      AtBat      Hits      HmRun      Runs      RBI
## 274.0200994 -0.3521900 -1.6371383  5.8146692  1.5423361  1.1241837
##      Walks      Years    CAtBat    CHits    CHmRun    CRuns
##   3.7288406 -16.3795195 -0.6411235  3.1629444  3.4005281 -0.9739405
##      CRBI      CWalks    LeagueN  DivisionW    PutOuts    Assists
##  -0.6003976   0.3378422 119.1434637 -144.0853061  0.1976300  0.6804200
##      Errors  NewLeagueN
##  -4.7127879 -71.0898914
```

```
set.seed(1)
cv.out=cv.glmnet(x[train,],y[train],alpha=0)
plot(cv.out)
```



```
bestlam=cv.out$lambda.min
bestlam
```

```
## [1] 326.0828
```

```
ridge.pred=predict(ridge.mod,s=bestlam,newx=x[test,])
mean((ridge.pred-y.test)^2)
```

```
## [1] 139856.6
```


```
out=glmnet(x,y,alpha=0)
predict(out,type="coefficients",s=bestlam)[1:20,]
```

```
## (Intercept)      AtBat      Hits      HmRun      Runs      RBI
## 15.44383120  0.07715547  0.85911582  0.60103106  1.06369007  0.87936105
##      Walks      Years      CAtBat      CHits      CHmRun      CRuns
##  1.62444617  1.35254778  0.01134999  0.05746654  0.40680157  0.11456224
##      CRBI      CWalks      LeagueN      DivisionW      PutOuts      Assists
##  0.12116504  0.05299202  22.09143197 -79.04032656  0.16619903  0.02941950
##      Errors      NewLeagueN
## -1.36092945  9.12487765
```


The Lasso

```
lasso.mod=glmnet(x[train,],y[train],alpha=1,lambda=grid)
plot(lasso.mod)
```

```
## Warning in regularize.values(x, y, ties, missing(ties), na.rm = na.rm):
## collapsing to unique 'x' values
```



```
set.seed(1)
cv.out=cv.glmnet(x[train,],y[train],alpha=1)
plot(cv.out)
```



```
bestlam=cv.out$lambda.min
lasso.pred=predict(lasso.mod,s=bestlam,newx=x[test,])
mean((lasso.pred-y.test)^2)
```

```
## [1] 143673.6
```

```
out=glmnet(x,y,alpha=1,lambda=grid)
lasso.coef=predict(out,type="coefficients",s=bestlam)[1:20,]
lasso.coef
```

##	(Intercept)	AtBat	Hits	HmRun	Runs
##	1.27479059	-0.05497143	2.18034583	0.00000000	0.00000000
##	RBI	Walks	Years	CAtBat	CHits
##	0.00000000	2.29192406	-0.33806109	0.00000000	0.00000000
##	CHmRun	CRuns	CRBI	CWalks	LeagueN
##	0.02825013	0.21628385	0.41712537	0.00000000	20.28615023
##	DivisionW	PutOuts	Assists	Errors	NewLeagueN
##	-116.16755870	0.23752385	0.00000000	-0.85629148	0.00000000

```
lasso.coef[lasso.coef!=0]
```

```
##      (Intercept)      AtBat      Hits      Walks      Years
##      1.27479059    -0.05497143    2.18034583    2.29192406   -0.33806109
##      CHmRun      CRuns      CRBI      LeagueN      DivisionW
##      0.02825013    0.21628385    0.41712537    20.28615023  -116.16755870
##      PutOuts      Errors
##      0.23752385   -0.85629148
```

Chapter 6 Lab 3: PCR and PLS Regression

Principal Components Regression

```
library(pls)
```

```
## Warning: package 'pls' was built under R version 4.3.3
```

```
##
```

```
## Attaching package: 'pls'
```

```
## The following object is masked from 'package:stats':
```

```
##
```

```
##      loadings
```

```
set.seed(2)
```

```
pcr.fit=pcr(Salary~., data=Hitters,scale=TRUE,validation="CV")
```

```
summary(pcr.fit)
```

```
## Data:      X dimension: 263 19
```

```
## Y dimension: 263 1
```

```
## Fit method: svdpc
```

```
## Number of components considered: 19
```

```
##
```

```
## VALIDATION: RMSEP
```

```
## Cross-validated using 10 random segments.
```


```
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              452    351.9    353.2    355.0    352.8    348.4    343.6
## adjCV           452    351.6    352.7    354.4    352.1    347.6    342.7
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV       345.5    347.7    349.6    351.4    352.1    353.5    358.2
## adjCV     344.7    346.7    348.5    350.1    350.7    352.0    356.5
##      14 comps 15 comps 16 comps 17 comps 18 comps 19 comps
## CV       349.7    349.4    339.9    341.6    339.2    339.6
## adjCV     348.0    347.7    338.2    339.7    337.2    337.6
```

```
##
```

```
## TRAINING: % variance explained
```


##	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps
## X	38.31	60.16	70.84	79.03	84.29	88.63	92.26	94.96
## Salary	40.63	41.58	42.17	43.22	44.90	46.48	46.69	46.75
##	9 comps	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps	
## X	96.28	97.26	97.98	98.65	99.15	99.47	99.75	
## Salary	46.86	47.76	47.82	47.85	48.10	50.40	50.55	
##	16 comps	17 comps	18 comps	19 comps				
## X	99.89	99.97	99.99	100.00				
## Salary	53.01	53.85	54.61	54.61				

```
validationplot(pcr.fit, val.type="MSEP")
```



```
set.seed(1)
pcr.fit=pcr(Salary~., data=Hitters,subset=train,scale=TRUE, validation="CV")
validationplot(pcr.fit, val.type="MSEP")
```


Salary



```
pcr.pred=predict(pcr.fit,x[test,],ncomp=7)
mean((pcr.pred-y.test)^2)
```

```
## [1] 140751.3
```

```
pcr.fit=pcr(y~x,scale=TRUE,ncomp=7)
summary(pcr.fit)
```

```
## Data:      X dimension: 263 19
## Y dimension: 263 1
## Fit method: svdpc
## Number of components considered: 7
## TRAINING: % variance explained
##   1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X   38.31   60.16   70.84   79.03   84.29   88.63   92.26
## y   40.63   41.58   42.17   43.22   44.90   46.48   46.69
```

Partial Least Squares


```
set.seed(1)
pls.fit=plsr(Salary~., data=Hitters,subset=train,scale=TRUE, validation="CV")
summary(pls.fit)
```

```

## Data:      X dimension: 131 19
## Y dimension: 131 1
## Fit method: kernelppls
## Number of components considered: 19
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              428.3   325.5   329.9   328.8   339.0   338.9   340.1
## adjCV           428.3   325.0   328.2   327.2   336.6   336.1   336.6
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          339.0   347.1   346.4   343.4   341.5   345.4   356.4
## adjCV       336.2   343.4   342.8   340.2   338.3   341.8   351.1
##      14 comps 15 comps 16 comps 17 comps 18 comps 19 comps
## CV          348.4   349.1   350.0   344.2   344.5   345.0
## adjCV       344.2   345.0   345.9   340.4   340.6   341.1
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X           39.13   48.80   60.09   75.07   78.58   81.12   88.21   90.71
## Salary      46.36   50.72   52.23   53.03   54.07   54.77   55.05   55.66
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X           93.17   96.05   97.08   97.61   97.97   98.70   99.12
## Salary      55.95   56.12   56.47   56.68   57.37   57.76   58.08
##      16 comps 17 comps 18 comps 19 comps
## X           99.61   99.70   99.95   100.00
## Salary      58.17   58.49   58.56   58.62

```

```
validationplot(pls.fit, val.type="MSEP")
```



```
pls.pred=predict(pls.fit,x[test,],ncomp=2)
mean((pls.pred-y.test)^2)
```

```
## [1] 145367.7
```

```
pls.fit=plsr(Salary~., data=Hitters,scale=TRUE,ncomp=2)
summary(pls.fit)
```

```
## Data:      X dimension: 263 19
## Y dimension: 263 1
## Fit method: kernelpls
## Number of components considered: 2
## TRAINING: % variance explained
##           1 comps  2 comps
## X           38.08   51.03
## Salary      43.05   46.40
```